

MEMORANDUM: PRODUCT ROADMAP H2 2026

To: Executive Leadership Committee | **From:** VP of Product, Maya Lin

Subject: Strategic AI Platform Milestones & Enterprise Commitments | **Date:** May 18, 2026

Executive Summary

In H2 2026, DocuMind AI will expand from a single-tenant document search engine into a collaborative multimodal intelligence workbench. Our top strategic goal is reaching \$4.5M ARR by Q4 2026 while maintaining 99.9% uptime SLA.

Key Deliverables & Target Quarters

- **Q3 Milestone (Release v2.0):** Deep Hybrid Retrieval combining BM25 lexical keyword matching with dense vector representations from Gemini-Embedding-2.
- **Q3 Milestone (Release v2.1):** Real-time collaborative multi-user chat sessions with granular cursor synchronization.
- **Q4 Milestone (Release v3.0):** Automated cross-document synthesis generating 15-page comparative executive briefings in under 30 seconds.

Resource Allocations & Budgeting

Total approved capital allocation for H2 infrastructure is \$680,000, broken down as follows:

Budget Category	Allocation (\$)	Primary Vendor	Expected Outcome
GPU Compute (Inference)	\$320,000	Lambda Labs / RunPod	Support 50k concurrent requests
Managed Vector Storage	\$140,000	Qdrant Cloud Enterprise	Host 2.5 billion document vectors
Security & Pen-testing	\$90,000	Bishop Fox	SOC2 Type II recertification
Data Engineering Pipeline	\$130,000	Internal Tooling & AWS	Sub-2s ingestion on 100-page PDFs

Target Exit Velocity: Sustained Gross Margin above 78% by end of December 2026.